

1.

$$\mathbf{x} = \begin{pmatrix} \mathbf{x}_a \\ \mathbf{x}_b \\ \mathbf{x}_c \end{pmatrix}, \quad \boldsymbol{\mu} = \begin{pmatrix} \boldsymbol{\mu}_a \\ \boldsymbol{\mu}_b \\ \boldsymbol{\mu}_c \end{pmatrix}, \quad \Sigma = \begin{pmatrix} \Sigma_{aa} & \Sigma_{ab} & \Sigma_{ac} \\ \Sigma_{ba} & \Sigma_{bb} & \Sigma_{bc} \\ \Sigma_{ca} & \Sigma_{cb} & \Sigma_{cc} \end{pmatrix}, \quad \Sigma^{-1} = \Lambda = \begin{pmatrix} \Lambda_{aa} & \Lambda_{ab} & \Lambda_{ac} \\ \Lambda_{ba} & \Lambda_{bb} & \Lambda_{bc} \\ \Lambda_{ca} & \Lambda_{cb} & \Lambda_{cc} \end{pmatrix}.$$

$$\begin{pmatrix} \mathbf{x}_a \\ \mathbf{x}_b \\ \mathbf{x}_c \end{pmatrix} \sim \mathcal{N}\left(\begin{pmatrix} \boldsymbol{\mu}_a \\ \boldsymbol{\mu}_b \\ \boldsymbol{\mu}_c \end{pmatrix}, \Sigma\right).$$

Since \mathbf{x}_c is marginalized, $\Lambda_{ac}, \Lambda_{bc}, \Lambda_{ca}, \Lambda_{cb}, \Lambda_{cc}$ do not affect the marginal distribution in \mathbf{x}_a and \mathbf{x}_b . Thus,

$$\begin{pmatrix} \mathbf{x}_a \\ \mathbf{x}_b \end{pmatrix} \sim \mathcal{N}\left(\begin{pmatrix} \boldsymbol{\mu}_a \\ \boldsymbol{\mu}_b \end{pmatrix}, \begin{pmatrix} \Sigma_{aa} & \Sigma_{ab} \\ \Sigma_{ba} & \Sigma_{bb} \end{pmatrix}\right).$$

$$\mathcal{N}(\mathbf{x} | \boldsymbol{\mu}, \Sigma) = \frac{1}{(2\pi)^{D/2} |\Sigma|^{1/2}} \exp\left\{-\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^\top \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu})\right\}.$$

$$\begin{aligned} & -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^\top \Sigma^{-1} (\mathbf{x} - \boldsymbol{\mu}) \\ &= -\frac{1}{2} \begin{pmatrix} \mathbf{x}_a - \boldsymbol{\mu}_a \\ \mathbf{x}_b - \boldsymbol{\mu}_b \end{pmatrix}^\top \begin{pmatrix} \Lambda_{aa} & \Lambda_{ab} \\ \Lambda_{ba} & \Lambda_{bb} \end{pmatrix} \begin{pmatrix} \mathbf{x}_a - \boldsymbol{\mu}_a \\ \mathbf{x}_b - \boldsymbol{\mu}_b \end{pmatrix} \\ &= -\frac{1}{2} (\mathbf{x}_a - \boldsymbol{\mu}_a)^\top \Lambda_{aa} (\mathbf{x}_a - \boldsymbol{\mu}_a) - \frac{1}{2} (\mathbf{x}_a - \boldsymbol{\mu}_a)^\top \Lambda_{ab} (\mathbf{x}_b - \boldsymbol{\mu}_b) \\ &\quad - \frac{1}{2} (\mathbf{x}_b - \boldsymbol{\mu}_b)^\top \Lambda_{ba} (\mathbf{x}_a - \boldsymbol{\mu}_a) - \frac{1}{2} (\mathbf{x}_b - \boldsymbol{\mu}_b)^\top \Lambda_{bb} (\mathbf{x}_b - \boldsymbol{\mu}_b) \\ &= -\frac{1}{2} \left[(\mathbf{x}_a - \boldsymbol{\mu}_a)^\top \Lambda_{aa} (\mathbf{x}_a - \boldsymbol{\mu}_a) + (\mathbf{x}_b - \boldsymbol{\mu}_b)^\top \Lambda_{bb} (\mathbf{x}_b - \boldsymbol{\mu}_b) \right. \\ &\quad \left. + (\mathbf{x}_a - \boldsymbol{\mu}_a)^\top \Lambda_{ab} (\mathbf{x}_b - \boldsymbol{\mu}_b) + (\mathbf{x}_b - \boldsymbol{\mu}_b)^\top \Lambda_{ba} (\mathbf{x}_a - \boldsymbol{\mu}_a) \right] \\ &= -\frac{1}{2} \left[(\mathbf{x}_a - \boldsymbol{\mu}_a)^\top \Lambda_{aa} (\mathbf{x}_a - \boldsymbol{\mu}_a) + (\mathbf{x}_b - \boldsymbol{\mu}_b)^\top \Lambda_{bb} (\mathbf{x}_b - \boldsymbol{\mu}_b) \right. \\ &\quad \left. + 2(\mathbf{x}_a - \boldsymbol{\mu}_a)^\top \Lambda_{ab} (\mathbf{x}_b - \boldsymbol{\mu}_b) \right] \\ &= -\frac{1}{2} (\mathbf{x}_a - \boldsymbol{\mu}_a)^\top \Lambda_{aa} (\mathbf{x}_a - \boldsymbol{\mu}_a) - \frac{1}{2} (\mathbf{x}_b - \boldsymbol{\mu}_b)^\top \Lambda_{bb} (\mathbf{x}_b - \boldsymbol{\mu}_b) \\ &\quad - (\mathbf{x}_a - \boldsymbol{\mu}_a)^\top \Lambda_{ab} (\mathbf{x}_b - \boldsymbol{\mu}_b). \end{aligned}$$

Let

$$\mathbf{a} = \mathbf{x}_a - \boldsymbol{\mu}_a \quad \text{and} \quad \mathbf{b} = \mathbf{x}_b - \boldsymbol{\mu}_b.$$

$$\begin{aligned}
& -\frac{1}{2} (\mathbf{x}_a - \boldsymbol{\mu}_a)^\top \Lambda_{aa} (\mathbf{x}_a - \boldsymbol{\mu}_a) - \frac{1}{2} (\mathbf{x}_b - \boldsymbol{\mu}_b)^\top \Lambda_{bb} (\mathbf{x}_b - \boldsymbol{\mu}_b) \\
& - (\mathbf{x}_a - \boldsymbol{\mu}_a)^\top \Lambda_{ab} (\mathbf{x}_b - \boldsymbol{\mu}_b) \\
& = -\frac{1}{2} \mathbf{a}^\top \Lambda_{aa} \mathbf{a} - \frac{1}{2} \mathbf{b}^\top \Lambda_{bb} \mathbf{b} - \mathbf{a}^\top \Lambda_{ab} \mathbf{b} \\
& = -\frac{1}{2} \mathbf{a}^\top \Lambda_{aa} \mathbf{a} - \mathbf{a}^\top \Lambda_{ab} \mathbf{b} \quad (\frac{1}{2} \mathbf{b}^\top \Lambda_{bb} \mathbf{b} \text{ is independent of } \mathbf{a}) \\
& = -\frac{1}{2} (\mathbf{x}_a - \boldsymbol{\mu}_a)^\top \Lambda_{aa} (\mathbf{x}_a - \boldsymbol{\mu}_a) - (\mathbf{x}_a - \boldsymbol{\mu}_a)^\top \Lambda_{ab} \mathbf{b} \\
& = -\frac{1}{2} \mathbf{x}_a^\top \Lambda_{aa} \mathbf{x}_a + \mathbf{x}_a^\top \Lambda_{aa} \boldsymbol{\mu}_a - \frac{1}{2} \boldsymbol{\mu}_a^\top \Lambda_{aa} \boldsymbol{\mu}_a - \mathbf{x}_a^\top \Lambda_{ab} \mathbf{b} + \boldsymbol{\mu}_a^\top \Lambda_{ab} \mathbf{b} \\
& = -\frac{1}{2} \mathbf{x}_a^\top \Lambda_{aa} \mathbf{x}_a + \mathbf{x}_a^\top [\Lambda_{aa} \boldsymbol{\mu}_a - \Lambda_{ab} \mathbf{b}] - \frac{1}{2} \boldsymbol{\mu}_a^\top \Lambda_{aa} \boldsymbol{\mu}_a + \boldsymbol{\mu}_a^\top \Lambda_{ab} \mathbf{b}. \longrightarrow (1)
\end{aligned}$$

$$-\frac{1}{2} (\mathbf{x}_a - \boldsymbol{\mu})^\top \Lambda_{aa} (\mathbf{x}_a - \boldsymbol{\mu}) = -\frac{1}{2} \mathbf{x}_a^\top \Lambda_{aa} \mathbf{x}_a + \mathbf{x}_a^\top \Lambda_{aa} \boldsymbol{\mu} - \frac{1}{2} \boldsymbol{\mu}^\top \Lambda_{aa} \boldsymbol{\mu}. \longrightarrow (2)$$

From (1) and (2), comparing the term $\mathbf{x}_a^\top \Lambda_{aa} \boldsymbol{\mu}$ with $\mathbf{x}_a^\top [\Lambda_{aa} \boldsymbol{\mu}_a - \Lambda_{ab} \mathbf{b}]$ leads to:

$$\begin{aligned}
\Lambda_{aa} \boldsymbol{\mu} &= \Lambda_{aa} \boldsymbol{\mu}_a - \Lambda_{ab} \mathbf{b}, \\
\boldsymbol{\mu} &= \Lambda_{aa}^{-1} (\Lambda_{aa} \boldsymbol{\mu}_a - \Lambda_{ab} \mathbf{b}) = \boldsymbol{\mu}_a - \Lambda_{aa}^{-1} \Lambda_{ab} \mathbf{b}.
\end{aligned}$$

Substituting $\mathbf{b} = \mathbf{x}_b - \boldsymbol{\mu}_b$ gives the conditional mean:

$$\boldsymbol{\mu}_{a|b} = \boldsymbol{\mu}_a - \Lambda_{aa}^{-1} \Lambda_{ab} (\mathbf{x}_b - \boldsymbol{\mu}_b).$$

Focusing on the quadratic term in \mathbf{x}_a :

$$-\frac{1}{2} (\mathbf{x}_a - \boldsymbol{\mu}_a)^\top \Lambda_{aa} (\mathbf{x}_a - \boldsymbol{\mu}_a)$$

shows that Λ_{aa} is the precision w.r.t \mathbf{x}_a , thus

$$\Sigma_{a|b}^{-1} = \Lambda_{aa} \implies \Sigma_{a|b} = \Lambda_{aa}^{-1}.$$

Therefore, the conditional distribution is

$$p(\mathbf{x}_a | \mathbf{x}_b) = \mathcal{N}(\boldsymbol{\mu}_a - \Lambda_{aa}^{-1} \Lambda_{ab} (\mathbf{x}_b - \boldsymbol{\mu}_b), \Lambda_{aa}^{-1}).$$

2.

If $(A + B C D)^{-1} = A^{-1} - A^{-1} B (C^{-1} + D A^{-1} B)^{-1} D A^{-1}$, then by the definition of the inverse $A A^{-1} x = x$, for every vector x it must be that $(A + B C D) [A^{-1} - A^{-1} B (C^{-1} + D A^{-1} B)^{-1} D A^{-1}] x = x$

$$\text{Let } y = [A^{-1} - A^{-1} B (C^{-1} + D A^{-1} B)^{-1} D A^{-1}] x$$

$$\begin{aligned}
y &= [A^{-1} - A^{-1} B (C^{-1} + D A^{-1} B)^{-1} D A^{-1}] x \\
&= A^{-1} x - A^{-1} B (C^{-1} + D A^{-1} B)^{-1} D A^{-1} x \\
&= (A + B C D) \cdot y \\
&= (A + B C D) \cdot [A^{-1} x - A^{-1} B (C^{-1} + D A^{-1} B)^{-1} D A^{-1} x] \\
&= (A + B C D) [A^{-1} x] - (A + B C D) [A^{-1} B (C^{-1} + D A^{-1} B)^{-1} D A^{-1} x]
\end{aligned}$$

$$\begin{aligned}
& (A + B C D) \left[A^{-1} x \right] \\
&= A A^{-1} x + B C D \left[A^{-1} x \right] \\
&= x + B C D A^{-1} x
\end{aligned}$$

$$\begin{aligned}
& (A + B C D) \left[A^{-1} B (C^{-1} + D A^{-1} B)^{-1} D A^{-1} x \right] \\
&= \left[A A^{-1} B + B C D A^{-1} B \right] (C^{-1} + D A^{-1} B)^{-1} D A^{-1} x \\
&= \left[B + B C D A^{-1} B \right] (C^{-1} + D A^{-1} B)^{-1} D A^{-1} x \\
&= B \left[1 + C D A^{-1} B \right] (C^{-1} + D A^{-1} B)^{-1} D A^{-1} x \\
&= B \left[C C^{-1} + C D A^{-1} B \right] (C^{-1} + D A^{-1} B)^{-1} D A^{-1} x \\
&= B \left[C (C^{-1} + D A^{-1} B) \right] (C^{-1} + D A^{-1} B)^{-1} D A^{-1} x \\
&= B C \left[(C^{-1} + D A^{-1} B) (C^{-1} + D A^{-1} B)^{-1} \right] D A^{-1} x \\
&= B C D A^{-1} x
\end{aligned}$$

$$\begin{aligned}
& (A + B C D) \cdot y \\
&= \left[x + B C D A^{-1} x \right] - B C D A^{-1} x \\
&= x
\end{aligned}$$

As for every vector x ,

$$(A + B C D) \left[A^{-1} - A^{-1} B (C^{-1} + D A^{-1} B)^{-1} D A^{-1} \right] x = x,$$

and by definition

$$A A^{-1} x = x,$$

it follows that

$$A^{-1} - A^{-1} B (C^{-1} + D A^{-1} B)^{-1} D A^{-1} = (A + B C D)^{-1}$$

If \mathbf{A} is an $n \times n$ diagonal matrix, then computing \mathbf{A}^{-1} simply involves taking reciprocals of the diagonal entries. Even if n is very large, only n such operations are needed, eliminating the need for expensive factorizations or dense matrix inversions. As \mathbf{A}^{-1} is known, its computational cost can be considered negligible.

Let \mathbf{B} be an $n \times k$ matrix (with $n \gg k$) and \mathbf{C} be an invertible $k \times k$ matrix. Instead of inverting the full $n \times n$ matrix $\mathbf{A} + \mathbf{BCD}$, the Woodbury identity can be applied which shifts the inversion problem to a smaller $k \times k$ matrix:

$$\mathbf{C}^{-1} + \mathbf{D} \mathbf{A}^{-1} \mathbf{B}.$$

As k is small compared to n —for example 3 or 10—this smaller system 3×3 or 10×10 is much cheaper to invert in comparison to the $n \times n$ system, especially if n runs into the thousands or millions.

In terms of computational complexity, a naive inversion of $\mathbf{A} + \mathbf{BCD}$ costs $O(n^3)$ when $n \gg k$. With the Woodbury identity and a known \mathbf{A}^{-1} , the cost instead involves:

- Multiplications with \mathbf{B} and \mathbf{D} , which is on the order of $O(nk^2)$.
- The inversion of the $k \times k$ matrix \mathbf{C} , which costs $O(k^3)$.

Because k is smaller in comparison to n , this leads to a computational advantage and results in computing the right-hand side of the formula.

3.

$$\mathcal{N}(x \mid \mu, \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma^2} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

$$\mathcal{N}\left(x \mid \mu, \frac{1}{\lambda}\right) = \sqrt{\frac{\lambda}{2\pi}} \exp\left(-\frac{\lambda}{2}(x-\mu)^2\right)$$

$$\begin{aligned} p(\mathbf{X} \mid \mu, \lambda) &= \prod_{n=1}^N \sqrt{\frac{\lambda}{2\pi}} \exp\left(-\frac{\lambda}{2}(x_n - \mu)^2\right) \\ &= \left(\sqrt{\frac{\lambda}{2\pi}}\right)^N \exp\left(-\frac{\lambda}{2} \sum_{n=1}^N (x_n - \mu)^2\right) \\ &= \left(\frac{\lambda}{2\pi}\right)^{\frac{N}{2}} \exp\left(-\frac{\lambda}{2} \sum_{n=1}^N (x_n - \mu)^2\right) \\ &= \lambda^{\frac{N}{2}} (2\pi)^{-\frac{N}{2}} \exp\left(-\frac{\lambda}{2} \sum_{n=1}^N (x_n - \mu)^2\right). \end{aligned}$$

$$p(\mu, \lambda) = \mathcal{N}(\mu \mid \mu_0, (\beta \lambda)^{-1}) \text{Gam}(\lambda \mid a, b)$$

$$\mathcal{N}(\mu \mid \mu_0, (\beta \lambda)^{-1}) = \sqrt{\frac{\beta \lambda}{2\pi}} \exp\left(-\frac{\beta \lambda}{2}(\mu - \mu_0)^2\right)$$

$$\text{Gam}(\lambda \mid a, b) = \frac{b^a}{\Gamma(a)} \lambda^{a-1} \exp(-b\lambda)$$

$$p(\mu, \lambda) = \sqrt{\frac{\beta \lambda}{2\pi}} \exp\left(-\frac{\beta \lambda}{2}(\mu - \mu_0)^2\right) \cdot \frac{b^a}{\Gamma(a)} \lambda^{a-1} \exp(-b\lambda)$$

$$= \sqrt{\frac{\beta}{2\pi}} \lambda^{\frac{1}{2}} \cdot \frac{b^a}{\Gamma(a)} \lambda^{a-1} \cdot \exp\left(-\frac{\beta \lambda}{2}(\mu - \mu_0)^2 - b\lambda\right)$$

$$= \sqrt{\frac{\beta}{2\pi}} \frac{b^a}{\Gamma(a)} \lambda^{\frac{1}{2} + (a-1)} \exp\left(-\frac{\beta \lambda}{2}(\mu - \mu_0)^2 - b\lambda\right)$$

$$= \sqrt{\frac{\beta}{2\pi}} \frac{b^a}{\Gamma(a)} \lambda^{a-\frac{1}{2}} \exp\left(-\frac{\beta \lambda}{2}(\mu - \mu_0)^2 - b\lambda\right)$$

$$= \sqrt{\frac{\beta}{2\pi}} \frac{b^a}{\Gamma(a)} \lambda^{a-\frac{1}{2}} \exp\left(-\frac{\beta \lambda}{2}\mu^2 + \beta \lambda \mu \mu_0 - \frac{\beta \lambda}{2}\mu_0^2 - b\lambda\right)$$

$$p(\mu, \lambda) = \sqrt{\frac{\beta}{2\pi}} \frac{b^a}{\Gamma(a)} \lambda^{a-\frac{1}{2}} \exp\left(-\frac{\beta \lambda \mu^2}{2} + \beta \mu_0 \lambda \mu - \frac{\beta \lambda}{2} \mu_0^2 - b \lambda\right) \longrightarrow (1)$$

Equation (2.153):

$$\begin{aligned} p(\mu, \lambda) &\propto \left[\lambda^{\frac{1}{2}} \exp\left(-\frac{\lambda \mu^2}{2}\right)\right]^\beta \exp(c \lambda \mu - d \lambda) \\ &= \left[\lambda^{\frac{1}{2}}\right]^\beta \left[\exp\left(-\frac{\lambda \mu^2}{2}\right)\right]^\beta \exp(c \lambda \mu - d \lambda) \\ &= \lambda^{\frac{\beta}{2}} \exp\left(-\frac{\beta \lambda \mu^2}{2}\right) \exp(c \lambda \mu - d \lambda) \\ &= \lambda^{\frac{\beta}{2}} \exp\left(-\frac{\beta \lambda \mu^2}{2} + c \lambda \mu - d \lambda\right) \longrightarrow (2). \end{aligned}$$

From (1) and (2):

$$\text{From the power of } \lambda : a - \frac{1}{2} = \frac{\beta}{2} \implies a = \frac{\beta}{2} + \frac{1}{2} = \frac{\beta + 1}{2}.$$

$$\text{From the } \lambda \mu \text{ term: } \beta \mu_0 = c \implies \mu_0 = \frac{c}{\beta}.$$

$$\text{From the constant term in the exponent: } -\frac{\beta \mu_0^2}{2} - b = -d \implies d = b + \frac{\beta \mu_0^2}{2}.$$

Substituting $\mu_0 = \frac{c}{\beta}$ gives $\frac{\beta \mu_0^2}{2} = \frac{\beta}{2} \left(\frac{c}{\beta}\right)^2 = \frac{c^2}{2\beta}$, so that $b = d - \frac{c^2}{2\beta}$.

$$\lambda^{\beta/2} = \lambda^{a-\frac{1}{2}} \implies a - \frac{1}{2} = \frac{\beta}{2} \implies a = \frac{1+\beta}{2}$$

The parameters of the distribution are:

$$\mu_0 = \frac{c}{\beta}, \quad a = \frac{\beta + 1}{2}, \quad b = d - \frac{c^2}{2\beta}$$

$$\begin{aligned} p(\mu, \lambda | \mathbf{X}) &\propto p(\mathbf{X} | \mu, \lambda) \cdot p(\mu, \lambda) \\ &\propto \lambda^{\frac{N}{2}} \exp\left(-\frac{\lambda}{2} \sum_{n=1}^N (x_n - \mu)^2\right) \cdot \lambda^{\frac{\beta}{2}} \exp\left(-\frac{\beta \lambda \mu^2}{2} + c \lambda \mu - d \lambda\right) \\ &\propto \lambda^{\frac{N+\beta}{2}} \exp\left[-\frac{\lambda}{2} \left(\sum_{n=1}^N (x_n - \mu)^2 + \beta \mu^2\right) + c \lambda \mu - d \lambda\right] \\ &\propto \lambda^{\frac{N+\beta}{2}} \exp\left[-\frac{\lambda}{2} \left(\sum_{n=1}^N x_n^2 - 2\mu \sum_{n=1}^N x_n + N\mu^2 + \beta \mu^2\right) + c \lambda \mu - d \lambda\right] \\ &\propto \lambda^{\frac{N+\beta}{2}} \exp\left[-\frac{\lambda}{2} \sum_{n=1}^N x_n^2 - d \lambda + \lambda \mu \left(\sum_{n=1}^N x_n + c\right) - \frac{\lambda}{2} (N + \beta) \mu^2\right] \\ &\propto \lambda^{\frac{N+\beta}{2}} \exp\left[-\frac{\lambda(N+\beta)\mu^2}{2} + \left(c + \sum_{n=1}^N x_n\right) \lambda \mu - \left(d + \frac{1}{2} \sum_{n=1}^N x_n^2\right) \lambda\right] \end{aligned}$$

$$d' = d + \frac{1}{2} \sum_{n=1}^N x_n^2$$

$$\begin{aligned}\mu'_0 &= \frac{c'}{\beta'} = \frac{c + \sum_{n=1}^N x_n}{N + \beta} \\ a' &= \frac{\beta + N + 1}{2} \\ b' &= d' - \frac{c'^2}{2\beta'}\end{aligned}$$

By matching exponents and powers of λ and μ , the posterior distribution is the same functional form as the prior distribution but with updated parameters. Under the Normal–Gamma parameterization:

$$p(\mu, \lambda | \mathbf{X}) = \mathcal{N}(\mu | \mu'_0, (\beta' \lambda)^{-1}) \Gamma(\lambda | a', b'),$$

where

$$\beta' = \beta + N, \quad c' = c + \sum_{n=1}^N x_n, \quad \mu'_0 = \frac{c'}{\beta'} = \frac{c + \sum_{n=1}^N x_n}{N + \beta}, \quad a' = \frac{\beta + N + 1}{2} \quad b' = d' - \frac{c'^2}{2\beta'}.$$

The posterior distribution is also a Gaussian–Gamma distribution of the same functional form as the prior, but with updated parameters, confirming that the posterior remains a Gaussian–Gamma distribution of the same functional form as the prior.

4.

a.

Wishart as a conjugate prior to $\Lambda = \Sigma^{-1}$ for Gaussian distribution $\mathcal{N}(\mu, \Lambda^{-1})$

A Wishart distribution over the precision matrix $\Lambda \in \mathbb{R}^{D \times D}$ with parameters (W, ν) is given by:

$$\mathcal{W}(\Lambda | W, \nu) = B(W, \nu) |\Lambda|^{\frac{\nu-D-1}{2}} \exp\left[-\frac{1}{2} \text{Tr}(W^{-1}\Lambda)\right] \rightarrow (1)$$

$$\mathcal{W}(\Lambda | W, \nu) \propto |\Lambda|^{\frac{\nu-D-1}{2}} \exp\left[-\frac{1}{2} \text{Tr}(W^{-1}\Lambda)\right]$$

$$p(X | \mu, \Lambda) \propto |\Lambda|^{\frac{N}{2}} \exp\left[-\frac{1}{2} \sum_{n=1}^N (x_n - \mu)^T \Lambda (x_n - \mu)\right]$$

Let $S = \frac{1}{N} \sum_{n=1}^N (x_n - \mu)(x_n - \mu)^T$. Then $\sum_{n=1}^N (x_n - \mu)^T \Lambda (x_n - \mu) = N \text{Tr}(\Lambda S)$.

$$p(X | \mu, \Lambda) \propto |\Lambda|^{\frac{N}{2}} \exp\left[-\frac{1}{2} \text{Tr}(N S \Lambda)\right] = |\Lambda|^{\frac{N}{2}} \exp\left[-\frac{N}{2} \text{Tr}(\Lambda S)\right]$$

Posterior in Λ :

$$\begin{aligned}p(\Lambda | X, W, \nu) &\propto p(X | \mu, \Lambda) \mathcal{W}(\Lambda | W, \nu) \\ &= |\Lambda|^{\frac{N}{2}} \exp\left[-\frac{1}{2} \text{Tr}(N S \Lambda)\right] \times |\Lambda|^{\frac{\nu-D-1}{2}} \exp\left[-\frac{1}{2} \text{Tr}(W^{-1}\Lambda)\right] \\ &= |\Lambda|^{\frac{N}{2}} \times |\Lambda|^{\frac{\nu-D-1}{2}} \cdot \exp\left[-\frac{1}{2} \text{Tr}(N S \Lambda)\right] \exp\left[-\frac{1}{2} \text{Tr}(W^{-1}\Lambda)\right] \\ &= |\Lambda|^{\frac{N+\nu-D-1}{2}} \cdot \exp\left[-\frac{1}{2} \text{Tr}((W^{-1} + N S) \Lambda)\right] \rightarrow (2)\end{aligned}$$

From (1) and (2),

$$\mathcal{W}(\Lambda | (W^{-1} + N S)^{-1}, N + \nu) = |\Lambda|^{\frac{N+\nu-D-1}{2}} \exp\left(-\frac{1}{2} \text{Tr}((W^{-1} + N S) \Lambda)\right)$$

$$\mathcal{W}(\Lambda | W, \nu) = B(W, \nu) |\Lambda|^{\frac{\nu-D-1}{2}} \exp\left[-\frac{1}{2} \text{Tr}(W^{-1}\Lambda)\right]$$

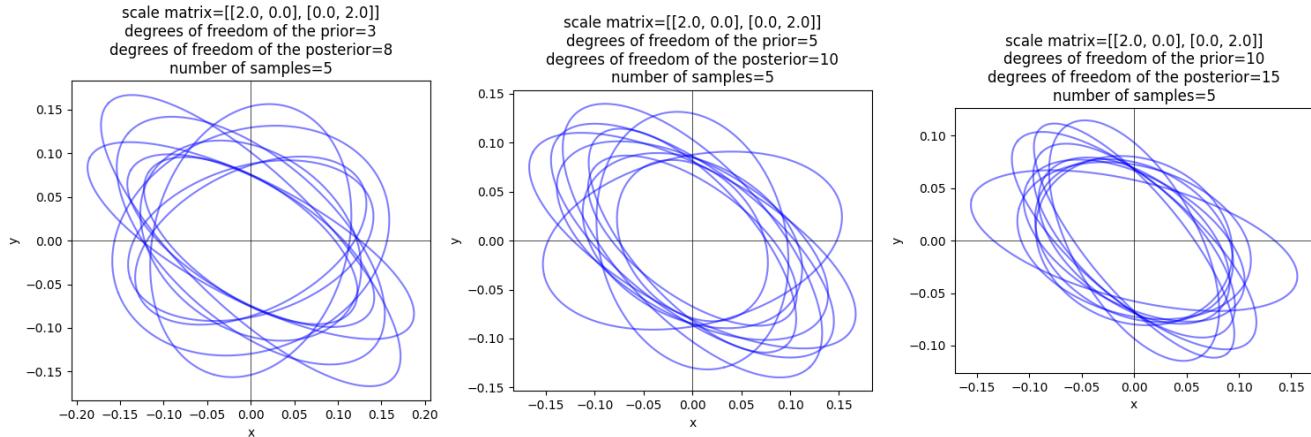
Therefore, the posterior is in the same form as a Wishart distribution with parameters:

$$\nu = N + \nu, \quad W = (W^{-1} + N S)$$

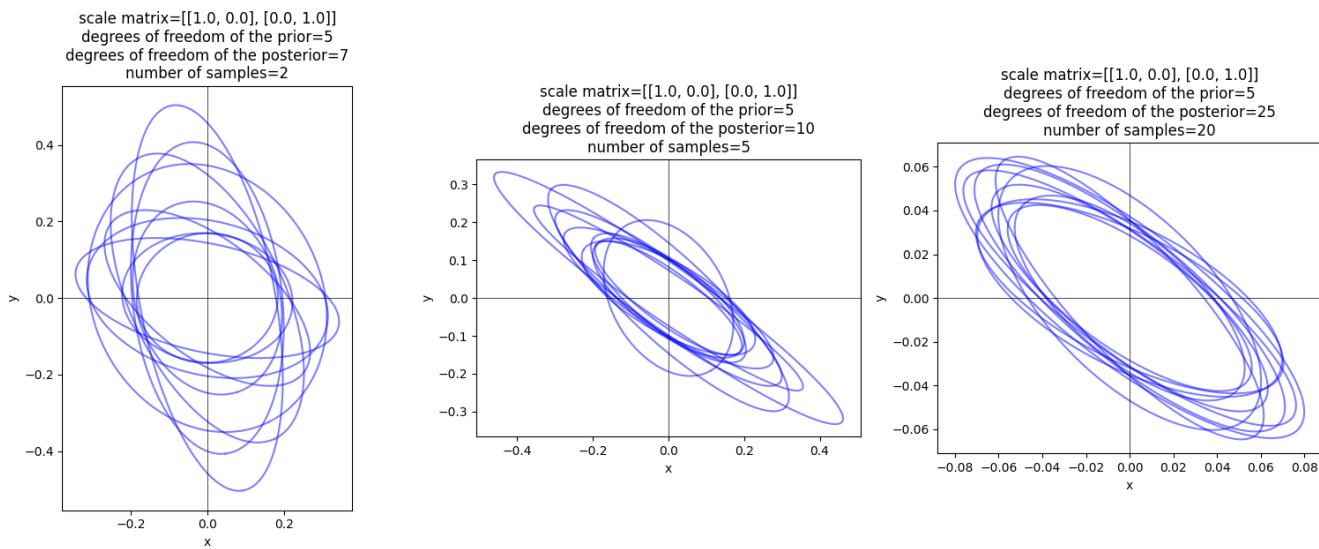
As $\mathcal{W}(\Lambda | W, \nu)$ (Wishart Prior) $\times p(X | \mu, \Lambda)$ (Gaussian likelihood) $= \mathcal{W}(\Lambda | W, \nu)$ (Wishart posterior), this closure property under posterior updates defines a conjugate prior for the Gaussian precision matrix.

b.

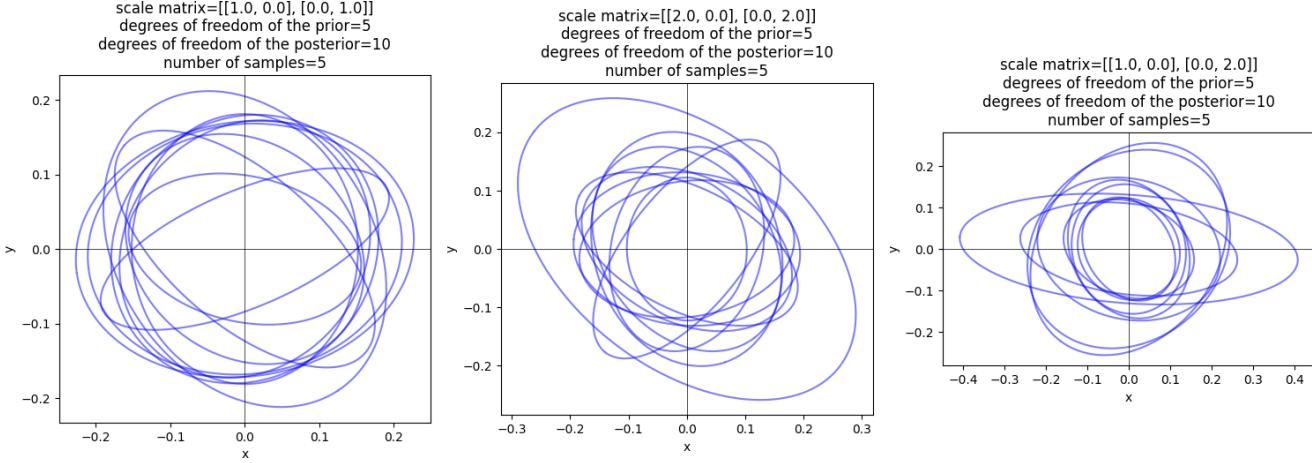
Varying Prior Degrees of Freedom



Varying Number of Samples



Varying Scale Matrix



```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 from scipy.stats import wishart
4
5 def generate_data(n, cov):
6     return np.random.multivariate_normal(mean=[0, 0], cov=cov, size=n)
7
8 def sample_wishart(df, scale, num_samples=5):
9     samples = []
10    for _ in range(num_samples):
11        W = wishart.rvs(df=df, scale=scale)
12        samples.append(W)
13    return samples
14
15 def plot_precision_ellipses(precisions, ax, title):
16    for p in precisions:
17        eigvals, eigvecs = np.linalg.eigh(p)
18        if np.any(eigvals <= 0):
19            continue
20
21        angles = np.linspace(0, 2*np.pi, 200)
22        circle = np.stack([np.cos(angles), np.sin(angles)], axis=1)
23        scale_matrix = np.diag(1.0 / np.sqrt(eigvals))
24        ellipse_y = circle @ scale_matrix
25        ellipse_x = ellipse_y @ eigvecs.T
26        ax.plot(ellipse_x[:, 0], ellipse_x[:, 1], 'b', alpha=0.5)
27
28    ax.set_aspect('equal', 'box')
29    ax.set_title(title)
30    ax.set_xlabel('x')
31    ax.set_ylabel('y')
32    ax.axhline(0, color='black', linewidth=0.5)
33    ax.axvline(0, color='black', linewidth=0.5)
34
35 def variation_1_distribution():
36     n = 5
37     scale_matrix = np.eye(2) * 2.0

```

```

38 cov = np.array([[1.0, 0.5], [0.5, 1.0]])
39 X = generate_data(n, cov)
40 sum_T = X.T @ X
41 dfs = [3, 5, 10]
42
43 fig, axes = plt.subplots(1, 3, figsize=(15, 4))
44
45 for i, df_prior in enumerate(dfs):
46     df_post = df_prior + n
47     scale_post = scale_matrix + sum_T
48     posterior_samples = sample_wishart(df_post, scale_post, num_samples=10)
49
50     title = (
51         f"scale_matrix={scale_matrix.tolist()}\n"
52         f"degrees_of_freedom_of_the_prior={df_prior}\n"
53         f"degrees_of_freedom_of_the_posterior={df_post}\n"
54         f"number_of_samples={n}"
55     )
56     plot_precision_ellipses(posterior_samples, axes[i], title)
57
58 fig.suptitle("Varying Prior Degrees of Freedom", fontsize=16)
59 plt.tight_layout()
60 plt.show()
61
62 def variation_2_distribution():
63     df_prior = 5
64     scale_matrix = np.eye(2)
65     cov = np.array([[1.0, 0.5], [0.5, 1.0]])
66     ns = [2, 5, 20]
67     fig, axes = plt.subplots(1, 3, figsize=(15, 4))
68
69     for i, n in enumerate(ns):
70         X = generate_data(n, cov)
71         sum_T = X.T @ X
72         df_post = df_prior + n
73         scale_post = scale_matrix + sum_T
74         posterior_samples = sample_wishart(df_post, scale_post, num_samples=10)
75
76         title = (
77             f"scale_matrix={scale_matrix.tolist()}\n"
78             f"degrees_of_freedom_of_the_prior={df_prior}\n"
79             f"degrees_of_freedom_of_the_posterior={df_post}\n"
80             f"number_of_samples={n}"
81     )
82     plot_precision_ellipses(posterior_samples, axes[i], title)
83
84 fig.suptitle("Varying Number of Samples", fontsize=16)
85 plt.tight_layout()
86 plt.show()
87
88 def variation_3_distribution():
89     df_prior = 5
90     n = 5
91     scale_matrix_list = [np.eye(2), 2.0 * np.eye(2), np.diag([1.0, 2.0])]
92     cov = np.array([[1.0, 0.5], [0.5, 1.0]])
93     X = generate_data(n, cov)
94     sum_T = X.T @ X
95
96     fig, axes = plt.subplots(1, 3, figsize=(15, 4))
97
98     for i, scale_matrix in enumerate(scale_matrix_list):
99         df_post = df_prior + n
100        scale_post = scale_matrix + sum_T
101        posterior_samples = sample_wishart(df_post, scale_post, num_samples=10)
102
103        title = (
104            f"scale_matrix={scale_matrix.tolist()}\n"
105            f"degrees_of_freedom_of_the_prior={df_prior}\n"
106            f"degrees_of_freedom_of_the_posterior={df_post}\n"
107            f"number_of_samples={n}"
108    )

```

```

109     plot_precision_ellipses(posterior_samples, axes[i], title)
110
111 fig.suptitle("Varying Scale Matrix", fontsize=16)
112 plt.tight_layout()
113 plt.show()
114
115 def main():
116     variation_1_distribution()
117     variation_2_distribution()
118     variation_3_distribution()
119
120 if __name__ == "__main__":
121     main()

```

5.

$$\begin{aligned}
N(\mathbf{x} | \boldsymbol{\mu}, \boldsymbol{\Sigma}) &= \frac{1}{(2\pi)^{\frac{D}{2}} |\boldsymbol{\Sigma}|^{\frac{1}{2}}} \exp\left\{-\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})\right\} \\
&= \frac{1}{(2\pi)^{\frac{D}{2}} |\boldsymbol{\Sigma}|^{\frac{1}{2}}} \exp\left\{-\frac{1}{2} \left(\mathbf{x}^\top \boldsymbol{\Sigma}^{-1} \mathbf{x} - 2\boldsymbol{\mu}^\top \boldsymbol{\Sigma}^{-1} \mathbf{x} + \boldsymbol{\mu}^\top \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu} \right)\right\} \\
&= \frac{1}{(2\pi)^{\frac{D}{2}} |\boldsymbol{\Sigma}|^{\frac{1}{2}}} \exp\left\{-\frac{1}{2} \mathbf{x}^\top \boldsymbol{\Sigma}^{-1} \mathbf{x} + \boldsymbol{\mu}^\top \boldsymbol{\Sigma}^{-1} \mathbf{x} - \frac{1}{2} \boldsymbol{\mu}^\top \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu}\right\} \\
&= \frac{1}{(2\pi)^{\frac{D}{2}} |\boldsymbol{\Sigma}|^{\frac{1}{2}}} \exp\left\{-\frac{1}{2} \mathbf{x}^\top \boldsymbol{\Sigma}^{-1} \mathbf{x} + \boldsymbol{\mu}^\top \boldsymbol{\Sigma}^{-1} \mathbf{x}\right\} \exp\left\{-\frac{1}{2} \boldsymbol{\mu}^\top \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu}\right\} \\
&= (2\pi)^{-\frac{D}{2}} \cdot |\boldsymbol{\Sigma}|^{-\frac{1}{2}} \exp\left\{-\frac{1}{2} \boldsymbol{\mu}^\top \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu}\right\} \exp\left\{-\frac{1}{2} \mathbf{x}^\top \boldsymbol{\Sigma}^{-1} \mathbf{x} + \boldsymbol{\mu}^\top \boldsymbol{\Sigma}^{-1} \mathbf{x}\right\}
\end{aligned}$$

Exponential-family form:

$$p(\mathbf{x} | \boldsymbol{\eta}) = h(\mathbf{x}) \cdot g(\boldsymbol{\eta}) \cdot \exp(\boldsymbol{\eta}^\top \mathbf{u}(\mathbf{x}))$$

$$h(\mathbf{x}) = (2\pi)^{-\frac{D}{2}}$$

$$\mathbf{u}(\mathbf{x}) = (\mathbf{x}, \mathbf{x}\mathbf{x}^\top) \quad (\text{which generalizes the univariate } (x, x^2))$$

$$\boldsymbol{\eta} = \begin{pmatrix} \boldsymbol{\eta}_1 \\ \boldsymbol{\eta}_2 \end{pmatrix} = \begin{pmatrix} \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu} \\ -\frac{1}{2} \boldsymbol{\Sigma}^{-1} \end{pmatrix}$$

Solving for $\boldsymbol{\Sigma}$ and $\boldsymbol{\mu}$ in terms of $\boldsymbol{\eta}$.

$$\boldsymbol{\eta}_2 = -\frac{1}{2} \boldsymbol{\Sigma}^{-1} \implies \boldsymbol{\Sigma}^{-1} = -2 \boldsymbol{\eta}_2 \implies \boldsymbol{\Sigma} = (\boldsymbol{\Sigma}^{-1})^{-1} = (-2 \boldsymbol{\eta}_2)^{-1} = -\frac{1}{2} \boldsymbol{\eta}_2^{-1}.$$

$$\boldsymbol{\eta}_1 = \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu} \implies \boldsymbol{\mu} = \boldsymbol{\Sigma} \boldsymbol{\eta}_1 = \left(-\frac{1}{2} \boldsymbol{\eta}_2^{-1}\right) \boldsymbol{\eta}_1$$

$$|\boldsymbol{\Sigma}|^{-\frac{1}{2}} = |-2 \boldsymbol{\eta}_2|^{\frac{D}{2}} \quad (\text{for } D \text{ dimensions})$$

$$g(\boldsymbol{\eta}) = |\boldsymbol{\Sigma}|^{-\frac{1}{2}} \exp\left(-\frac{1}{2} \boldsymbol{\mu}^\top \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu}\right)$$

$$\boldsymbol{\mu}^\top \Sigma^{-1} \boldsymbol{\mu} = (\Sigma \boldsymbol{\eta}_1)^\top \Sigma^{-1} (\Sigma \boldsymbol{\eta}_1) = \boldsymbol{\eta}_1^\top \Sigma \boldsymbol{\eta}_1 = \boldsymbol{\eta}_1^\top \left(-\frac{1}{2} \boldsymbol{\eta}_2^{-1} \right) \boldsymbol{\eta}_1 = -\frac{1}{2} \boldsymbol{\eta}_1^\top \boldsymbol{\eta}_2^{-1} \boldsymbol{\eta}_1$$

Therefore,

$$\begin{aligned} g(\boldsymbol{\eta}) &= |\Sigma|^{-\frac{1}{2}} \exp\left(-\frac{1}{2} \boldsymbol{\mu}^\top \Sigma^{-1} \boldsymbol{\mu}\right) \\ &= |\Sigma|^{-\frac{1}{2}} \exp\left(-\frac{1}{2} \boldsymbol{\eta}_1^\top \boldsymbol{\eta}_2^{-1} \boldsymbol{\eta}_1\right) \\ &= |-2 \boldsymbol{\eta}_2|^{\frac{D}{2}} \exp\left(\frac{1}{4} \boldsymbol{\eta}_1^\top \boldsymbol{\eta}_2^{-1} \boldsymbol{\eta}_1\right) \end{aligned}$$

Univariate ($D = 1$) case:

$$g(\boldsymbol{\eta}) = |-2 \eta_2|^{\frac{1}{2}} \exp\left(\frac{\eta_1^2}{4 \eta_2}\right)$$

Multivariate ($D > 1$) case:

$$g(\boldsymbol{\eta}) = |-2 \boldsymbol{\eta}_2|^{\frac{D}{2}} \exp\left(\frac{1}{4} \boldsymbol{\eta}_1^\top \boldsymbol{\eta}_2^{-1} \boldsymbol{\eta}_1\right)$$